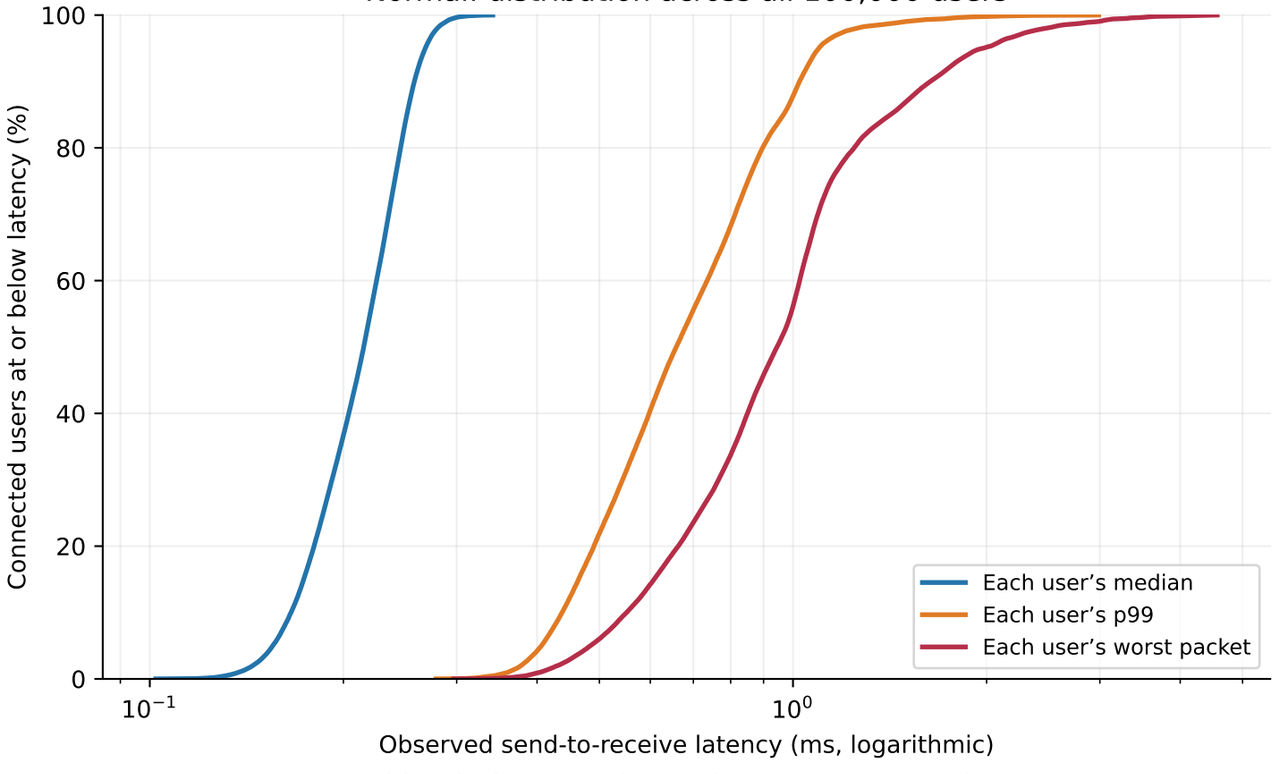
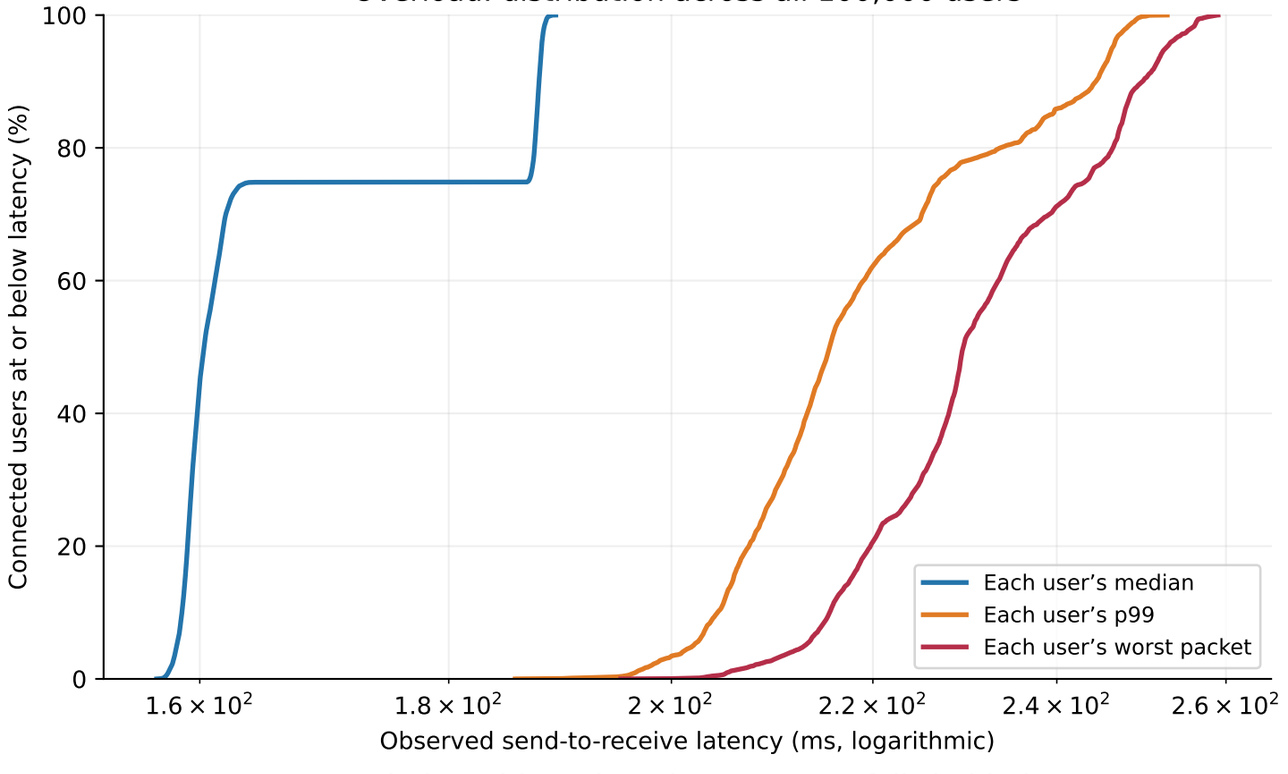
Realtime relay latency fairness · selected Mio runtime · 100,000 connected users

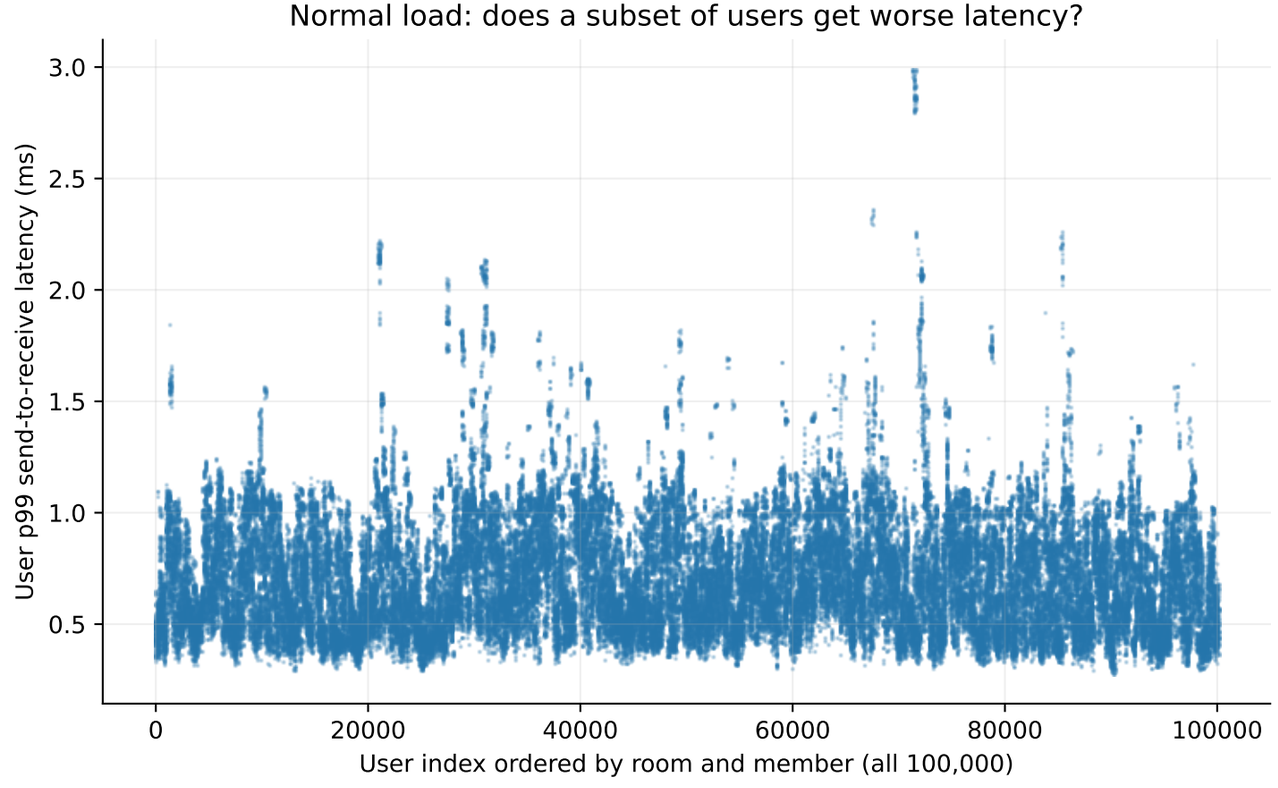
Normal: distribution across all 100,000 users



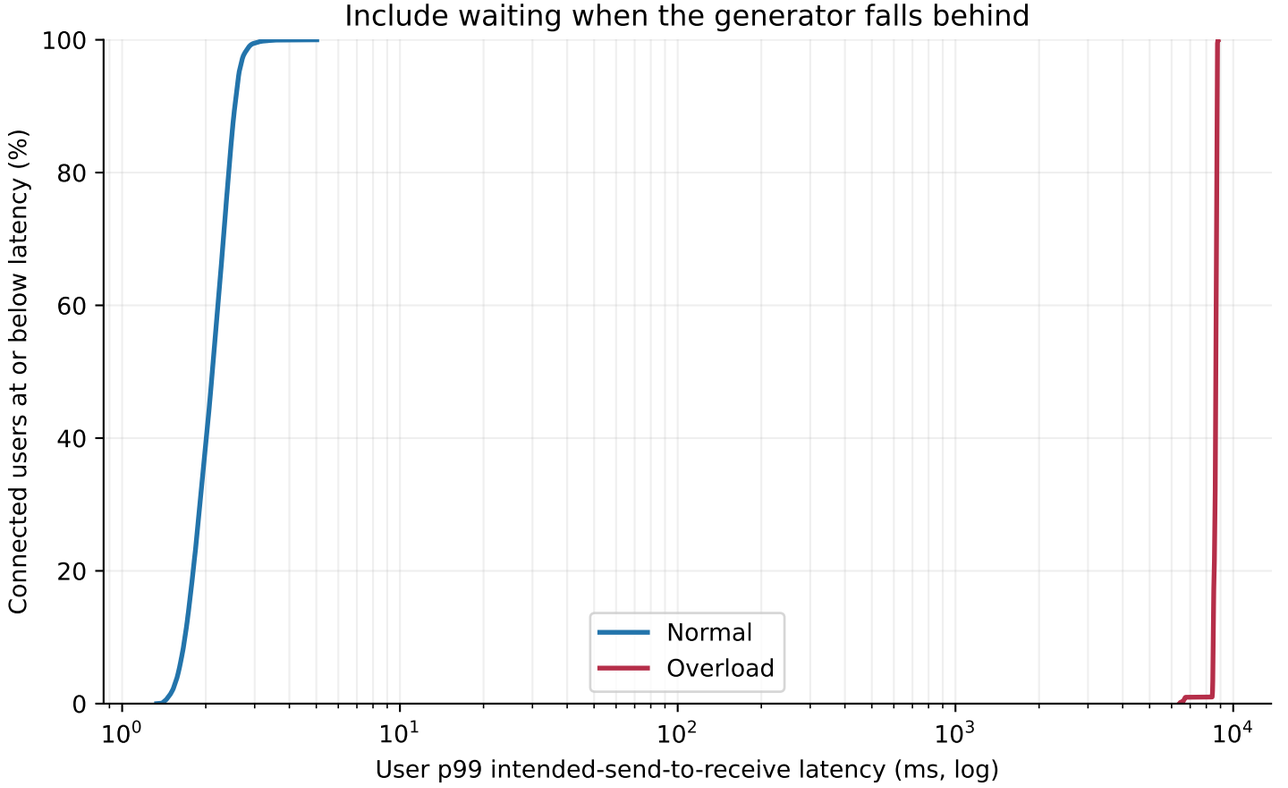
Overload: distribution across all 100,000 users



Normal load: does a subset of users get worse latency?



Include waiting when the generator falls behind



1 KiB private messages, 16-user rooms, 4 server workers / 8 client cores. Normal: 2 packets/user/s × 60 s; overload: 10 × 10 s.
Exact per-user samples; empirical tails from 120 / 100 packets each. Local loopback includes client scheduling and batched sends.